

1 MODULE *Consensus*

The consensus problem requires a set of processes to choose a single value. This module specifies the problem by specifying exactly what the requirements are for choosing a value.

7 EXTENDS *Naturals*, *FiniteSets*, *TLAPS*

We let the constant parameter *Value* be the set of all values that can be chosen.

13 CONSTANT *Value*

We now specify the safety property of consensus as a trivial algorithm that describes the allowed behaviors of a consensus algorithm. It uses the variable *chosen* to represent the set of all chosen values. The algorithm is trivial; it allows only behaviors that contain a single state-change in which the variable *chosen* is changed from its initial value $\{\}$ to the value $\{v\}$ for an arbitrary value v in *Value*. The algorithm itself does not specify any fairness properties, so it also allows a behavior in which *chosen* is not changed. We could use a translator option to have the translation include a fairness requirement, but we don't bother because it is easy enough to add it by hand to the safety specification that the translator produces.

A real specification of consensus would also include additional variables and actions. In particular, it would have Propose actions in which clients propose values and Learn actions in which clients learn what value has been chosen. It would allow only a proposed value to be chosen. However, the interesting part of a consensus algorithm is the choosing of a single value. We therefore restrict our attention to that aspect of consensus algorithms. In practice, given the algorithm for choosing a value, it is obvious how to implement the Propose and Learn actions.

For convenience, we define the macro *Choose()* that describes the action of changing the value of *chosen* from $\{\}$ to $\{v\}$, for a nondeterministically chosen v in the set *Value*. (There is little reason to encapsulate such a simple action in a macro; however our other specs are easier to read when written with such macros, so we start using them now.) The *when* statement can be executed only when its condition, *chosen* = $\{\}$, is true. Hence, at most one *Choose()* action can be performed in any execution. The *with* statement executes its body for a nondeterministically chosen v in *Value*. Execution of this statement is enabled only if *Value* is non-empty—something we do not assume at this point because it is not required for the safety part of consensus, which is satisfied if no value is chosen.

We put the *Choose()* action inside a *while* statement that loops forever. Of course, only a single *Choose()* action can be executed. The algorithm stops after executing a *Choose()* action. Technically, the algorithm deadlocks after executing a *Choose()* action because control is at a statement whose execution is never enabled. Formally, termination is simply deadlock that we want to happen. We could just as well have omitted the *while* and let the algorithm terminate. However, adding the *while* loop makes the TLA+ representation of the algorithm a tiny bit simpler.

```

62 --algorithm Consensus{
63   variable chosen =  $\{\}$ ;
64   macro Choose( ) { when chosen =  $\{\}$ ;
65                       with (  $v \in \textit{Value}$  ) { chosen :=  $\{v\}$  } }
66   { lbl: while ( TRUE ) { Choose() }
67   }
68 }
```

The *PlusCal* translator writes the TLA+ translation of this algorithm below. The formula *Spec* is the TLA+ specification described by the algorithm's code. For now, you should just understand its two subformulas *Init* and *Next*. Formula *Init* is the initial predicate and describes all possible initial states of an execution. Formula *Next* is the next-state relation; it describes the possible state changes (changes of the values of variables), where unprimed variables represent their values in the old state and primed variables represent their values in the new state.

```

80  **** BEGIN TRANSLATION
81  VARIABLE chosen
82
83  vars  $\triangleq \langle \textit{chosen} \rangle$ 
84
85  Init  $\triangleq$  Global variables
86              $\wedge \textit{chosen} = \{\}$ 
87
88  Next  $\triangleq$   $\wedge \textit{chosen} = \{\}$ 
89              $\wedge \exists v \in \textit{Value} :$ 
90              $\textit{chosen}' = \{v\}$ 
91
92  Spec  $\triangleq \textit{Init} \wedge \Box[\textit{Next}]_{\textit{vars}}$ 
93
94  **** END TRANSLATION
95 |-----|
    We now prove the safety property that at most one value is chosen. We first define the type-
    correctness invariant TypeOK, and then define Inv to be the inductive invariant that asserts
    TypeOK and that the cardinality of the set chosen is at most 1. We then prove that, in any
    behavior satisfying the safety specification Spec, the invariant Inv is true in all states. This means
    that at most one value is chosen in any behavior.
105 TypeOK  $\triangleq$   $\wedge \textit{chosen} \subseteq \textit{Value}$ 
106              $\wedge \textit{IsFiniteSet}(\textit{chosen})$ 
107
108 Inv  $\triangleq$   $\wedge \textit{TypeOK}$ 
109              $\wedge \textit{Cardinality}(\textit{chosen}) \leq 1$ 
110
    To prove our theorem, we need the following simple results about the cardinality of finite sets.
115 AXIOM EmptySetCardinality  $\triangleq \textit{IsFiniteSet}(\{\}) \wedge \textit{Cardinality}(\{\}) = 0$ 
116 AXIOM SingletonCardinality  $\triangleq$ 
117              $\forall e : \textit{IsFiniteSet}(\{e\}) \wedge (\textit{Cardinality}(\{e\}) = 1)$ 
118
    Whenever we add an axiom, we should check it with TLC to make sure we haven't made any
    errors. To check axiom SingletonCardinality, we must replace the unbounded quantification with
    a bounded one. We therefore let TLC check that the following formula is true.
125 SingleCardinalityTest  $\triangleq$ 
126              $\forall e \in \text{SUBSET } \{\text{"a"}, \text{"b"}, \text{"c"}\} : \textit{IsFiniteSet}(\{e\}) \wedge (\textit{Cardinality}(\{e\}) = 1)$ 
127
    We now prove that Inv is an invariant, meaning that it is true in every state in every behavior.
    Before trying to prove it, we should first use TLC to check that it is true. It's hardly worth
    bothering to either check or prove the obvious fact that Inv is an invariant, but it's a nice tiny
    exercise. Model checking is instantaneous when Value is set to any small finite set.
    To understand the following proof, you need to understand the formula Spec, which equals
     $\textit{Init} \wedge \Box[\textit{Next}]_{\textit{vars}}$ 
    where vars is the tuple  $\langle \textit{chosen}, \textit{pc} \rangle$  of all variables. It is a temporal formula satisfied by a
    behavior iff the behavior starts in a state satisfying Init and such that each step (sequence of
    states) satisfies  $[\textit{Next}]_{\textit{vars}}$ , which equals

```

$Next \vee (vars' = vars)$

Thus, each step satisfies either *Next* (so it is a step allowed by the next-state relation) or it is a “stuttering step” that leaves all the variables unchanged. The reason why a spec must allow stuttering steps will become apparent when we prove that a consensus algorithm satisfies this specification of consensus.

By default, a definition is not usable in a proof. If the a definition should be usable in the proof of a step, meaning that the prover can expand the definition, then that must be explicitly indicated—usually in the `DEF` clause of a `BY` statement. A `DEF` clause of a `USE` statement makes the definitions in it usable in the scope of that statement. The following `USE` statement makes the definition of `lbl` usable everywhere in the rest of the module. (There is a corresponding `HIDE` statement that makes definitions unusable in its scope.)

165 `USE DEF lbl`

The following lemma asserts that *Inv* is an inductive invariant of the next-state action *Next*. It is the key step in proving that *Inv* is an invariant of (true in every behavior allowed by) specification *Spec*.

```

172 LEMMA InductiveInvariance  $\triangleq$ 
173        $Inv \wedge [Next]_{vars} \Rightarrow Inv'$ 
174 <1>1. SUFFICES ASSUME Inv,  $[Next]_{vars}$ 
175       PROVE Inv'
176   OBVIOUS
177 <1>2.CASE Next
178   <2>1. PICK  $v \in Value : chosen' = \{v\}$ 
179     In the following BY proof, <1>2 denotes the case assumption Next
180   BY <1>2 DEF Next
181   <2>2. QED
182     In the following BY proof, <1>1 refers to the assumptions
183     Inv and  $[Next]_{vars}$ 
184   BY <1>1, <2>1, SingletonCardinality,  $1 \leq 1$  DEF Inv, TypeOK, Next, vars
185 <1>4.CASE  $vars' = vars$ 
186   BY <1>1, <1>4 DEF Inv, TypeOK, Next, vars
187 <1>5. QED
188   BY <1>1, <1>2, <1>4 DEF Next

```

TLAPS does not yet handle temporal logic reasoning. Therefore, proofs of temporal steps are omitted. However, we indicate in comments what we expect the proofs to look like when *TLAPS* does prove temporal formulas.

```

196 THEOREM Invariance  $\triangleq Spec \Rightarrow \Box Inv$ 
197 <1>1. Init  $\Rightarrow Inv$ 
198   We usually have to use SimpleArithmetic to prove facts about numbers,
199   but  $0 \leq 1$  is simple enough that Isabelle can prove it by itself.
200   BY EmptySetCardinality,  $0 \leq 1$  DEF Init, Inv, TypeOK

202 <1>2. QED
203   <2>1.  $Inv \wedge \Box [Next]_{vars} \Rightarrow \Box Inv$  Typo corrected here 28 Jun 2016
204   BY InductiveInvariance, RuleINV1

```

205 PROOF OMITTED

RuleInv1 is defined in the *TLAPS* module by

THEOREM *RuleINV1* \triangleq ASSUME STATE *I*, STATE *F*, ACTION *N*,
 $I \wedge [N]_F \Rightarrow I'$
 PROVE $I \wedge \Box[N]_F \Rightarrow \Box I$

213 $\langle 2 \rangle 2$. QED

214 BY $\langle 1 \rangle 1$, $\langle 2 \rangle 1$

215 PROOF OMITTED

We now define *LiveSpec* to be the algorithm's specification with the added fairness condition of weak fairness of the next-state relation, which asserts that execution does not stop if some action is enabled. The temporal formula *Success* asserts that some value is eventually chosen. Below, we prove that *LiveSpec* implies *Success*. This means that, in every behavior satisfying *LiveSpec*, some value is chosen.

226 *LiveSpec* \triangleq *Spec* \wedge $WF_{vars}(Next)$

227 *Success* \triangleq $\Diamond(chosen \neq \{\})$

For liveness, we need to assume that there exists at least one value.

232 ASSUME *ValueNonempty* \triangleq *Value* $\neq \{\}$

TLAPS does not yet reason about *ENABLED*. Therefore, we must omit all proofs that involve *ENABLED* formulas. To perform as much of the proof as possible, as much as possible we restrict the use of an *ENABLED* expression to a step asserting that it equals its definition. *ENABLED A* is true of a state *s* iff there is a state *t* such that the step $s \rightarrow t$ satisfies *A*. It follows from this semantic definition that *ENABLED A* equals the formula obtained by

1. Expanding all definitions of defined symbols in *A* until all primes are priming variables.
2. For each primed variable, replacing every instance of that primed variable by a new symbol (the same symbol for each primed variable).
3. Existentially quantifying over those new symbols.

252 LEMMA *EnabledDef* \triangleq

253 *TypeOK* \Rightarrow

254 $((\text{ENABLED } \langle Next \rangle_{vars}) \equiv (chosen = \{\}))$

255 $\langle 1 \rangle$ DEFINE *E* \triangleq

256 $\exists chosenp :$

257 $\wedge \wedge chosen = \{\}$

258 $\wedge \exists v \in Value : chosenp = \{v\}$

259 $\wedge \neg(\langle chosenp \rangle = \langle chosen \rangle)$

260 $\langle 1 \rangle 1$. *E* = *ENABLED* $\langle Next \rangle_{vars}$

261 BY DEF *Next*, *vars*

262 PROOF OMITTED

263 $\langle 1 \rangle 2$. SUFFICES ASSUME *TypeOK*

264 PROVE *E* = $(chosen = \{\})$

265 BY $\langle 1 \rangle 1$

266 PROOF OMITTED

267 $\langle 1 \rangle 3$. *E* = $\exists chosenp : E!(chosenp)!1$

```

268 BY <1>2 DEF TypeOK
269 <1>4. @ = (chosen = {})
270 BY <1>2, ValueNonempty DEF TypeOK
271 <1>5. QED
272 BY <1>3, <1>4

```

Here is our proof that *Livespec* implies *Success*. It uses the standard TLA proof rules. For example *RuleWF1* is defined in the *TLAPS* module to be the rule *WF1* discussed in

```

AUTHOR = "Leslie Lamport",
TITLE  = "The Temporal Logic of Actions",
JOURNAL = topas,
volume = 16,
number = 3,
YEAR   = 1994,
month  = may,
PAGES  = "872--923"

```

PTL stands for propositional temporal logic reasoning. We expect that, when *TLAPS* handles temporal reasoning, it will use a decision procedure for *PTL*.

```

292 THEOREM LiveSpec ⇒ Success
293 <1>1. □Inv ∧ □[Next]vars ∧ WFvars(Next) ⇒ (chosen = {} ∼ chosen ≠ {})
294 <2>1. SUFFICES □[Next]vars ∧ WFvars(Next) ⇒ ((Inv ∧ chosen = {}) ∼ chosen ≠ {})
295     OBVIOUS
296     PROOF OMITTED
297 <2>2. (Inv ∧ (chosen = {})) ∧ [Next]vars ⇒ ((Inv' ∧ (chosen' = {})) ∨ chosen' ≠ {})
298     BY InductiveInvariance
299 <2>3. (Inv ∧ (chosen = {})) ∧ <Next>vars ⇒ (chosen' ≠ {})
300     BY DEF Inv, Next, vars
301 <2>4. (Inv ∧ (chosen = {})) ⇒ ENABLED <Next>vars
302     BY EnabledDef DEF Inv
303     PROOF OMITTED
304 <2>5. QED
305     BY <2>2, <2>3, <2>4, RuleWF1
306     PROOF OMITTED
307 <1>2. (chosen = {} ∼ chosen ≠ {}) ⇒ ((chosen = {}) ⇒ ◇(chosen ≠ {}))
308     OBVIOUS
309     PROOF OMITTED
310 <1>3. QED
311     BY Invariance, <1>1, <1>2 DEF LiveSpec, Spec, Init, Success
312     PROOF OMITTED

```

The following theorem is used in the refinement proof in module *VoteProof*.

```

318 THEOREM LiveSpecEquals ≜
319     LiveSpec ≡ Spec ∧ (□◇<Next>vars ∨ □◇(chosen ≠ {}))
320 <1>1. ∧ Spec ≡ Spec ∧ □TypeOK
321     ∧ LiveSpec ≡ LiveSpec ∧ □TypeOK
322     BY Invariance DEF LiveSpec, Inv

```

```

323   PROOF OMITTED
324  ⟨1⟩2. □ TypeOK ⇒ ((□◇¬ENABLED ⟨Next⟩vars) ≡ □◇(chosen ≠ {}))
325  ⟨2⟩1. □ TypeOK ⇒ □((¬ENABLED ⟨Next⟩vars) ≡ (chosen ≠ {}))
326    BY EnabledDef
327   PROOF OMITTED
328  ⟨2⟩2. QED
329    BY ⟨2⟩1
330   PROOF OMITTED
331  ⟨1⟩3. QED
332    BY ⟨1⟩1, ⟨1⟩2 DEF LiveSpec
333   PROOF OMITTED
334  ┌──────────────────────────────────────────────────────────────────────────────────┐

```

```

\ * Modification History
\ * Last modified Sat Nov 16 22:17:07 CST 2019 by hengxin
\ * Last modified Tue Feb 14 13:35:49 PST 2012 by lamport
\ * Last modified Mon Feb 07 14:46:59 PST 2011 by lamport

```